

updates.sql

```
/*
 * Run a series of updates on the data
 */

/* Set the status as 'complete' for three low level tickets
 * These three low-level tickets were partitioned from the same high level
ticket
 * This is done on purpose, and will play into the next (more interesting)
update...
 */
update ticket
    set status = 'complete'
    where tid in (169, 170, 171);

-- -- -- -- --
-- -- -- -- --

/*
 * Update the status for all high-level tickets to 'complete',
 * iff all the low-level tickets partitioned from said high level tickets
 * have status 'complete'.
 *
 * After the previous update, all the low-level tickets associated with the
 * high level ticket with ticket id `1` (associated low-level tickets have
ticket id's
 * 169, 170, 171), will be set to 'complete', but none others will be
altered.
 */
with ticketstatuses as (
    select ht.tid,
           t.status,
           lt.tid as lowtickettid,
           tt.status as lowticketstatus
    from highticket ht
    inner join ticket t
        on ht.tid = t.tid
    inner join lowticket lt
        on ht.tid = lt.highticketid
    inner join ticket tt
        on lt.tid = tt.tid
    order by lowtickettid
```

updates.sql

```
)
update ticket
  set status = 'complete'
  where tid in (
    select distinct t.tid
    from(
      select *,
        count(1) over(partition by tid) as numlowleveltickets,
        sum(case
          when lowticketstatus = 'complete'
            then 1
            else 0
          end) over(partition by tid) as
numcompletelowleveltickets
      from ticketstatuses
    ) t
    where t.numlowleveltickets = t.numcompletelowleveltickets
  );
```

```
-- --
-- --

/*
 * It might be the case that an employee is on sick/maternity leave,
 * and all of his/her responsibilities need to be transferred over to another
employee.
 *
 * The employee on leave may be a developer - in which case low-level
tickets may be assigned to him/her -
 * or a business owner - in which case he/she may be accountable for one or
more projects.
 *
 * Below, I'll transfer all the responsibilities of employee with person id
'2' to the employee with person id '1'.
 *
 * The trouble with this is that the responsibilities of an employee may span
across multiple tables -
 * for example, a ticket may be assigned to that employee, that employee may
be an accountable business owner.
 * This adds complexity because the column that needs to be changed in each
of the tables may be different -
```

updates.sql

```
* for example in the table `ticketdeveloperassignment` the id of the
developer to which a ticket is assigned is called `devid`,
* which should be changed, while in the table `projectassignment`, the
column which needs to be changed is `pid`.
*
* To do this all in one shot, I'll write an code block (DO statement) to be
run database side.
* This block will read a temporary table, which specifies the tables to
change and the columns to change for each table,
* and change the specified columns with value `2` to `1` - transferring all
responsibilites from person with pid `2` to
* the person with pid `1`.
*/
-- create a temporary table which specifies which tables to change,
-- and which columns within each table to change.
drop table if exists temp_tablestochange;
create temporary table temp_tablestochange(
    tbl varchar(50),
    coltochange varchar(50)
);
-- insert the tables/columns to change into this temporary table.
insert into temp_tablestochange values
('ticketdeveloperassignment', 'devid'),
('accountableowner', 'pid'),
('projectassignment', 'pid');

-- Do statement - change the specified tables.
-- This works by formatting the string 'UPDATE <tbl> SET <col> = 1 where
<col> = 2' - and executing it as SQL.
DO $$DECLARE r record;
BEGIN
    FOR r IN SELECT tbl, coltochange from temp_tablestochange
    LOOP
        EXECUTE 'UPDATE ' || r.tbl || ' set ' || r.coltochange || ' = 1 where '
|| r.coltochange || ' = 2';
    END LOOP;
END$$;

--
--
/*
```

updates.sql

```
* Overly ambitious developers/project managers may assign tickets to a
sprint
* on the last day of the sprint's lifespan.
*
* It's not a reasonable thing to do, however, as it's not possible to
finish an entire
* ticket in one day, and so the ticket should probably be kept saved until
the next sprint.
*
* These commands isolate only the tickets that were scheduled to a sprint
at least
* one day before the sprint's closing date, and stores it in a temporary
table.
* It then removes all the values in the `ticketsprintassignment` table, and
replaces them
* with the values in the temporary table.
*
* This effectively cleans the `ticketsprintassignment` of values that leave
less than one
* day to complete a ticket.
*/
-- create a temporary table of the ticket assignments that were assigned
-- to a sprint at least one day before the sprint's closing
create temporary table temp_reasonableticketschedules as (
    with assignmentinfo as (
        select tsa.tid,
            tsa.sprinnumber,
            tsa.teamname,
            tsa.datescheduled,
            s.startdate,
            s.enddate
        from ticketsprintassignment tsa
        inner join sprint s
            on tsa.sprinnumber = s.sprinnumber
            and tsa.teamname = s.teamname
        order by tid
    )
    select t.tid,
        t.sprinnumber,
        t.teamname,
        t.datescheduled
```

updates.sql

```
from
(
    select *,
        enddate - datescheduled as scheduledworktime
    from assignmentinfo
) t
where scheduledworktime > 0
);
-- delete all values from the `ticketsprintassignment` table.
delete from ticketsprintassignment
where true;
-- Replace these values with the values of the temporary table.
insert into ticketsprintassignment
select * from temp_reasonableticketschedules;
```